

Introduction

The dataset I have chosen contains a list of candy. The candy contains features such as chocolate, fruity, peanutyalmondy, nougat, crispedricewafer, hard, bar, pluribus and caramel. These are represented as binary variables, 1 means yes and 0 means no.

The question to be answered in this assignment is:

- 1) What is the most important feature?
- 2) Predicting if the candy will be chocolate?

Columns:

1. **chocolate**: Does the candy contain chocolate?
2. **fruity**: Is the candy fruity?
3. **caramel**: Does the candy contain caramel?
4. **peanutalmondy**: Does the candy contain almonds or peanuts?
5. **nougat**: Does it contain nougat
6. **crispedricewafer**: Does it contain crisped rice or wafers?
7. **hard**: Is it a hard candy?
8. **bar**: Is it a candy bar?
9. **pluribus**: Is it one of many candies in a bag or box?
10. **sugarpercent**: The percentile of sugar it falls under within the data set.
11. **pricepercent**: The unit price percentile compared to the rest of the set.
12. **winpercent**: The overall win percentage according to 269,000 matchups.

1) Imports

Pandas is a python library. It is used for data manipulation and analysis (Read_csv)

seaborn - data visualisation library. (Heatmap)

matplotlib.pyplot - embedding plots (Heatmap)

Sklearn tree - import decision tree (Tree)

Sklearn Linear_model - import linear model (Linear Model)

sklearn.model_selection import train_test_split - Creating Train and test set (Train&Test)

sklearn.linear_model import LogisticRegression - import logistic regression (LR)

sklearn import metrics - Metrics (accuracy)

sklearn.metrics import mean_squared_error - Test accuracy (Confusion Matrix)

```
In [17]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import tree
from sklearn import datasets, linear_model
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn import metrics
from sklearn.metrics import mean_squared_error, r2_score
```

2) Read in data

- The read_csv function reads in the dataset from github. read_csv can control delimiters, rows, column names. The data is stored in a dataframe candy
- drop_duplicates function removes any duplicates

```
In [18]: candy = pd.read_csv('https://raw.githubusercontent.com/fivethirtyeight/d
ata/master/candy-power-ranking/candy-data.csv')
candy.drop_duplicates(inplace=True)
```

3) Descriptive Statistics

Performing descriptive Statistics help understand the features of the candy data set by giving short summaries and measure of the data.

What I found:

- `.head()` displays first 5 rows in dataset.
- `.describe()` is used to view some basic statistical details like count, mean, std, min, 25%, 50%, 75% and max
- `.index` returns a list of the indexes
- `.columns` return all columns. There is a total of 13 columns
- `.dtypes` returns the data types of each column. Is used to check if data is categorical or numerical. All data is numerical except for competitorname. The numerical columns are divided into 9 integers and 3 floats
- `.shape` returns a tuple representing the dimensionality of the DataFrame. The result is (85,13) - 85 rows and 13 columns
- `.values` returns the values in the dataframe
- `.info()` prints information including the index dtype and columns, non-null values

In [19]: `candy.head()`

Out[19]:

	competitorname	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer	hard
0	100 Grand	1	0	1	0	0	1	0
1	3 Musketeers	1	0	0	0	1	0	0
2	One dime	0	0	0	0	0	0	0
3	One quarter	0	0	0	0	0	0	0
4	Air Heads	0	1	0	0	0	0	0

In [20]: `candy.describe()`

Out[20]:

	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer	hard
count	85.000000	85.000000	85.000000	85.000000	85.000000	85.000000	85.000000
mean	0.435294	0.447059	0.164706	0.164706	0.082353	0.082353	0.176471
std	0.498738	0.500140	0.373116	0.373116	0.276533	0.276533	0.383482
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
75%	1.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

In [21]: `candy.index`

Out[21]: Int64Index([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 75, 76, 77, 78, 79, 80, 81, 82, 83, 84],
dtype='int64')

In [6]: `candy.columns`

Out[6]: Index(['competitorname', 'chocolate', 'fruity', 'caramel', 'peanutyalmondy', 'nougat', 'crispedricewafer', 'hard', 'bar', 'pluribus', 'sugarpercent', 'pricepercent', 'winpercent'],
dtype='object')

```
In [22]: candy.dtypes
```

```
Out[22]: competitorname    object
         chocolate         int64
         fruity            int64
         caramel           int64
         peanutyalmondy    int64
         nougat            int64
         crispedricewafer  int64
         hard              int64
         bar               int64
         pluribus          int64
         sugarpercent      float64
         pricepercent      float64
         winpercent        float64
         dtype: object
```

```
In [23]: candy.shape
```

```
Out[23]: (85, 13)
```

```
In [24]: candy.values
```

```
Out[24]: array([[ '100 Grand', 1, 0, ..., 0.73199999, 0.86000001, 66.971725],
                [ '3 Musketeers', 1, 0, ..., 0.60399997, 0.51099998, 67.602936],
                [ 'One dime', 0, 0, ..., 0.011000000000000001, 0.11599999999999999
          9,
                32.261086],
                ...,
                [ 'WelchÕs Fruit Snacks', 0, 1, ..., 0.31299999, 0.31299999,
                44.375519],
                [ 'WertherÕs Original Caramel', 0, 0, ..., 0.18600000000000003,
                0.26699999, 41.904308],
                [ 'Whoppers', 1, 0, ..., 0.87199998, 0.84799999, 49.524113]],
          dtype=object)
```

```
In [25]: candy.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 85 entries, 0 to 84
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   competitorname        85 non-null    object
1   chocolate             85 non-null    int64
2   fruity                85 non-null    int64
3   caramel               85 non-null    int64
4   peanutyalmondy        85 non-null    int64
5   nougat                85 non-null    int64
6   crispedricewafer      85 non-null    int64
7   hard                  85 non-null    int64
8   bar                   85 non-null    int64
9   pluribus              85 non-null    int64
10  sugarpercent          85 non-null    float64
11  pricepercent           85 non-null    float64
12  winpercent             85 non-null    float64
dtypes: float64(3), int64(9), object(1)
memory usage: 9.3+ KB
```

3) More Descriptive Statistics

- `candy.sugarpercent.describe()` views the view some basic statistical details like count, mean, std, min, 25%, 50%, 75% and max on sugarpercent
- `candy.sugarpercent.mean()` will find the average percentage of sugar in each candy
- `candy.sugarpercent.max()` - max value of sugar percentage
- `candy.sugarpercent.min()` - min value of sugar percentage
- `candy[candy.sugarpercent > candy.sugarpercent.mean()].sort_values(by=['sugarpercent'], ascending=False)` - shows candies that are higher then the mean. sorts by top sugar percent. Top 5 are Reese's stuffed with pieces, Sugar Babies, Milky Way Simply Caramel and Skittles wildberry
- `candy[candy.sugarpercent <= candy.sugarpercent.mean()].sort_values(by=['sugarpercent'], ascending=True)` - shows candies that are lower then the mean. sorts by lowest sugar percent. The candies with the lowest sugar percentage are one dime, one quarter, Reese's Miniatures, Chiclets and lemonhead
- `candy[candy.sugarpercent==candy.sugarpercent.max()].competitorname` - Candy with the highest sugar content = Reese's stuffed with pieces
- `candy[candy['fruity']==0].sort_values(by=['winpercent'], ascending=False).head(10)` - sorts non fruity candies by winpercent. Top non fruity candies are Reese's Peanut Butter cup, Reese's Miniatures, Twix, Kit Kat and Snickers
- `candy[(candy['caramel']==1)&(candy['fruity']==1)]` - displays candy which is both fruity and has caramel = Caramel Apple Pops
- `sort = candy[['competitorname', 'winpercent']].sort_values(by='winpercent')`
`pd.concat([sort.head(5), sort.tail(5)], axis=0).plot(x='competitorname', y='winpercent popularity', sort_columns=True, figsize = (5,5))` - Data visualisation showing the most popular candies(Bar chart)

```
In [31]: candy.sugarpercent.describe()
```

```
Out[31]: count      85.000000
         mean        0.478647
         std         0.282778
         min         0.011000
         25%         0.220000
         50%         0.465000
         75%         0.732000
         max         0.988000
         Name: sugarpercent, dtype: float64
```

```
In [32]: candy.sugarpercent.mean()
```

```
Out[32]: 0.4786470514588237
```

```
In [33]: candy.sugarpercent.max()
```

```
Out[33]: 0.98799998
```

```
In [34]: candy.sugarpercent.min()
```

```
Out[34]: 0.011000000000000001
```



```
In [35]: candy[candy.sugarpercent > candy.sugarpercent.mean()].sort_values(by=['s  
ugarpercent'],ascending=False)
```

Out[35]:

	competitorname	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer	hard
54	ReeseOs stuffed with pieces	1	0	0	1	0	0	0
70	Sugar Babies	0	0	1	0	0	0	0
38	Milky Way Simply Caramel	1	0	1	0	0	0	0
61	Skittles wildberry	0	1	0	0	0	0	0
60	Skittles original	0	1	0	0	0	0	0
4	Air Heads	0	1	0	0	0	0	0
8	Candy Corn	0	0	0	0	0	0	0
17	Gobstopper	0	1	0	0	0	0	1
84	Whoppers	1	0	0	0	0	1	0
34	Mike & Ike	0	1	0	0	0	0	0
58	Runts	0	1	0	0	0	0	1
56	Rolo	1	0	1	0	0	0	0
41	Nerds	0	1	0	0	0	0	1
33	M&MOs	1	0	0	0	0	0	0
32	Peanut butter M&MOs	1	0	0	1	0	0	0
55	Ring pop	0	1	0	0	0	0	1
57	Root Beer Barrels	0	0	0	0	0	0	1
0	100 Grand	1	0	1	0	0	1	0
37	Milky Way Midnight	1	0	1	0	1	0	0
16	Fun Dip	0	1	0	0	0	0	1
13	Dots	0	1	0	0	0	0	0
14	Dum Dums	0	1	0	0	0	0	1
11	Chewey Lemonhead Fruit Mix	0	1	0	0	0	0	0
52	ReeseOs Peanut Butter cup	1	0	0	1	0	0	0
36	Milky Way	1	0	1	0	1	0	0
49	Pop Rocks	0	1	0	0	0	0	1
1	3 Musketeers	1	0	0	0	1	0	0
10	Charleston Chew	1	0	0	0	1	0	0
9	Caramel Apple Pops	0	1	1	0	0	0	0
65	Snickers Crisper	1	0	1	1	0	1	0

	competitorname	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer	hard
6	Baby Ruth	1	0	1	1	1	0	0
73	Swedish Fish	0	1	0	0	0	0	0
74	Tootsie Pop	1	1	0	0	0	0	1
42	Nestle Butterfinger	1	0	0	1	0	0	0
47	Peanut M&Ms	1	0	0	1	0	0	0
50	Red vines	0	1	0	0	0	0	0
69	Strawberry bon bons	0	1	0	0	0	0	1
64	Snickers	1	0	1	1	1	0	0
79	Twix	1	0	1	0	0	1	0

```
In [36]: candy[candy.sugarpercent <= candy.sugarpercent.mean() ].sort_values(by=[  
        'sugarpercent' ],ascending=True)
```

Out[36]:

	competitorname	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer	hard
2	One dime	0	0	0	0	0	0	0
3	One quarter	0	0	0	0	0	0	0
51	Reese's Miniatures	1	0	0	1	0	0	0
12	Chiclets	0	1	0	0	0	0	0
30	Lemonhead	0	1	0	0	0	0	1
67	Sour Patch Tricksters	0	1	0	0	0	0	0
66	Sour Patch Kids	0	1	0	0	0	0	0
81	Warheads	0	1	0	0	0	0	1
48	Pixie Sticks	0	0	0	0	0	0	0
26	Jawbusters	0	1	0	0	0	0	1
15	Fruit Chews	0	1	0	0	0	0	0
22	Hershey's Kisses	1	0	0	0	0	0	0
68	Starburst	0	1	0	0	0	0	0
72	Super Bubble	0	1	0	0	0	0	0
76	Tootsie Roll Midgies	1	0	0	0	0	0	0
83	Werther's Original Caramel	0	0	1	0	0	0	1
44	Nik L Nip	0	1	0	0	0	0	0
27	Junior Mints	1	0	0	0	0	0	0
59	Sixlets	1	0	0	0	0	0	0
29	Laffy Taffy	0	1	0	0	0	0	0
80	Twizzlers	0	1	0	0	0	0	0
45	Now & Later	0	1	0	0	0	0	0
31	Lifesavers big ring gummies	0	1	0	0	0	0	0
63	Smarties candy	0	1	0	0	0	0	1
62	Nestle Smarties	1	0	0	0	0	0	0
35	Milk Duds	1	0	1	0	0	0	0
78	Trolli Sour Bites	0	1	0	0	0	0	0
7	Boston Baked Beans	0	0	0	1	0	0	0
75	Tootsie Roll Juniors	1	0	0	0	0	0	0
28	Kit Kat	1	0	0	0	0	1	0

	competitorname	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer	hard
43	Nestle Crunch	1	0	0	0	0	1	0
82	Welch's Fruit Snacks	0	1	0	0	0	0	0
39	Mounds	1	0	0	0	0	0	0
40	Mr Good Bar	1	0	0	1	0	0	0
53	Reese's pieces	1	0	0	1	0	0	0
71	Sugar Daddy	0	0	1	0	0	0	0
23	Hershey's Krackel	1	0	0	0	0	1	0
24	Hershey's Milk Chocolate	1	0	0	0	0	0	0
25	Hershey's Special Dark	1	0	0	0	0	0	0
46	Payday	0	0	0	1	1	0	0
21	Haribo Twin Snakes	0	1	0	0	0	0	0
20	Haribo Sour Bears	0	1	0	0	0	0	0
19	Haribo Happy Cola	0	0	0	0	0	0	0
18	Haribo Gold Bears	0	1	0	0	0	0	0
77	Tootsie Roll Snack Bars	1	0	0	0	0	0	0
5	Almond Joy	1	0	0	1	0	0	0

```
In [37]: candy[candy.sugarpercent==candy.sugarpercent.max() ].competitorname
```

```
Out[37]: 54    Reese's stuffed with pieces
Name: competitorname, dtype: object
```

```
In [38]: candy[candy['fruity']==0].sort_values(by=['winpercent'], ascending=False)
         ).head(10)
```

Out[38]:

	competitorname	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer	hard
52	Reese's Peanut Butter cup	1	0	0	1	0	0	0
51	Reese's Miniatures	1	0	0	1	0	0	0
79	Twix	1	0	1	0	0	1	0
28	Kit Kat	1	0	0	0	0	1	0
64	Snickers	1	0	1	1	1	0	0
53	Reese's pieces	1	0	0	1	0	0	0
36	Milky Way	1	0	1	0	1	0	0
54	Reese's stuffed with pieces	1	0	0	1	0	0	0
32	Peanut butter M&M's	1	0	0	1	0	0	0
42	Nestle Butterfinger	1	0	0	1	0	0	0

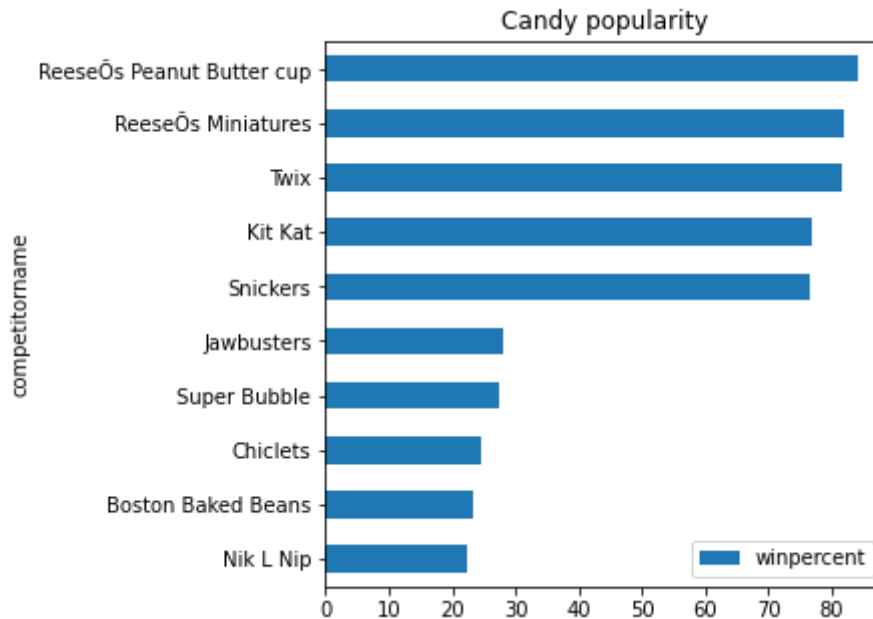
```
In [39]: candy[(candy['caramel']==1)&(candy['fruity']==1)]
```

Out[39]:

	competitorname	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer	hard
9	Caramel Apple Pops	0	1	1	0	0	0	0

```
In [40]: sort = candy[['competitorname', 'winpercent']].sort_values(by='winpercent')
pd.concat([sort.head(5), sort.tail(5)], axis=0).plot(x='competitorname', y='winpercent', kind='barh', title='Candy popularity', sort_columns=True, figsize = (5,5))
```

Out[40]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7582e45f60>



4) Exploratory Data Analysis

a) Correlation

Multivariate analysis

This is used to find the pairwise correlation of all the columns in the dataframe i.e. finds relationship between the features.

By looking at the correlation we can easily identify useful exploratory variables.

A negative pair number on the chart means the two candies are less likely both have those attributes

A positive pair number on the chart means the two candies are more likely both have those attributes

High Correlations between features normally results in poor linear and logistic Regression performance

In this plot we can see that chocolate candies are rarely fruity

Correlation lets us check for multicollinearity. The coefficients in this plot are normal and stable so we do not have to worry about this

A heat map is a visualisation technique where each value is depicted by colour. Basically it will show that get the most attention.

In my heatmap it ranges from dark blue to dark red. The more negative the value darker the shade of blue. The more positive the darker the shade of red.

This heatmap is used to make it even easier to identify the correlations


```
In [41]: Correlations=candy.corr()  
Correlations
```

Out[41]:

	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer	
chocolate	1.000000	-0.741721	0.249875	0.377824	0.254892	0.341210	-
fruity	-0.741721	1.000000	-0.335485	-0.399280	-0.269367	-0.269367	-
caramel	0.249875	-0.335485	1.000000	0.059356	0.328493	0.213113	-
peanutyalmondy	0.377824	-0.399280	0.059356	1.000000	0.213113	-0.017646	-
nougat	0.254892	-0.269367	0.328493	0.213113	1.000000	-0.089744	-
crispedricewafer	0.341210	-0.269367	0.213113	-0.017646	-0.089744	1.000000	-
hard	-0.344177	0.390678	-0.122355	-0.205557	-0.138675	-0.138675	-
bar	0.597421	-0.515066	0.333960	0.260420	0.522976	0.423751	-
pluribus	-0.339675	0.299725	-0.269585	-0.206109	-0.310339	-0.224693	-
sugarpercent	0.104169	-0.034393	0.221933	0.087889	0.123081	0.069950	-
pricepercent	0.504675	-0.430969	0.254327	0.309153	0.153196	0.328265	-
winpercent	0.636517	-0.380938	0.213416	0.406192	0.199375	0.324680	-

```
In [42]: plt.figure(figsize = (10,6))  
sns.heatmap(candy.corr(),annot=True, cmap = 'coolwarm')
```

Out[42]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7582e457f0>



B) Density Plots

Also known as Kernel Density Plots or Density Trace Graph. A Density Plot visualises the distribution of data over a continuous interval or time period. The peaks of these densities can show where the values are concentrated.

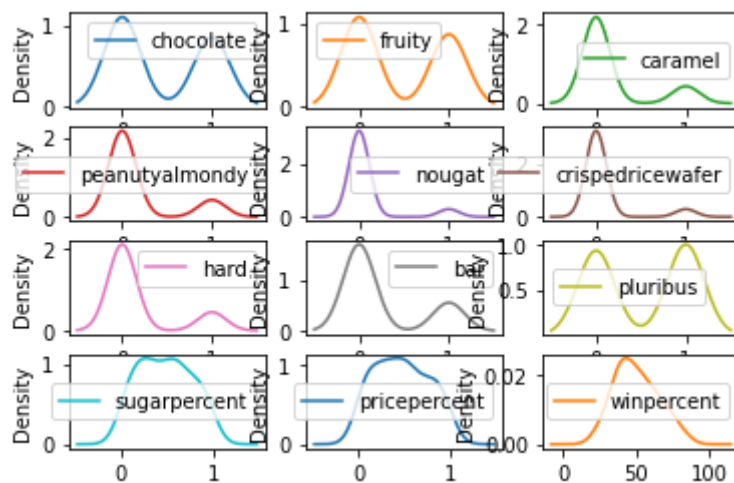
An advantage of density plots over histograms is that it is better at determining the distribution shape. This is because it is not affected by the bins.

By plotting a density plot we change the kind to 'density'

From this plot we can see:

- More candies have no caramel
- more candies are chocolatey
- Nearly all candies are not hard

```
In [43]: K=candy.plot(kind='density',subplots=True,layout=(4,3),sharex=False,sharey=False)
```



C) Box plots

Univariate Analysis

Boxplots summarize the distribution of each attribute, drawing a line for the median. They graphically depict groups of numerical data through their quartiles

Gives an idea of data spread

The winpercent.quantile is set to 60% and above. So if the relationship is above this value. The feature can be considered important.

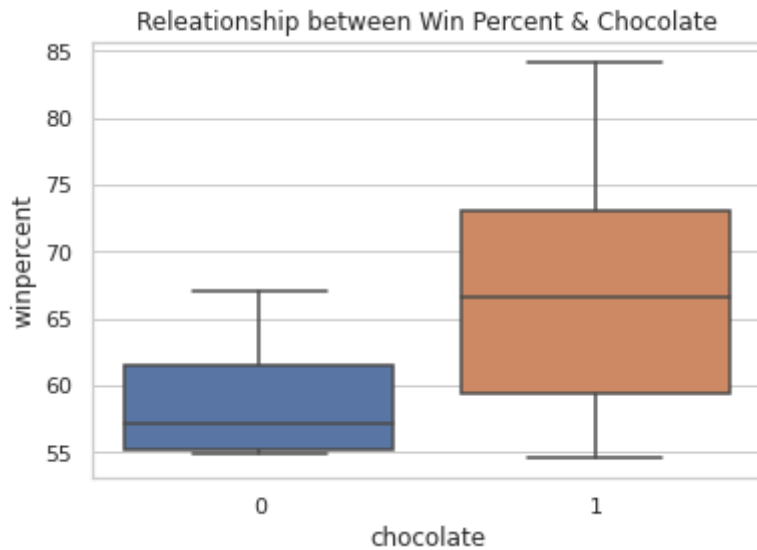
Results of boxplots:

- Win Percent & Chocolate = <60% like chocolate candies, Therefore liking chocolate might be an important parameter
- Win Percent & Fruity = <60% like fruity candies, Therefore liking fruity candies might be an important parameter
- Win Percent & Caramel = <60% like caramel candies, Therefore liking caramel candies might be an important parameter
- Win Percent & PeanutAlmond = <60% like peanutalmond candies, Therefore liking peanutalmond candies might be an important parameter
- Win Percent & Nougat = <60% like Nougat candies, Therefore liking peanutalmond candies might be an important parameter
- Win Percent & Wafer = <60% like wafer candies, Therefore liking Wafer candies might be an important parameter
- Win Percent & Hardness = <60% like soft candies, Therefore Soft Candies might be important
- Win Percent & Bar = <60% are bars, Therefore being a bar might be an important parameter
- Win Percent & Pluribus = <60% , Not clear result


```
In [44]: win_num = candy[candy.winpercent>candy.winpercent.quantile(.6)]
```

```
In [45]: sns.set_theme(style="whitegrid")
sns.boxplot(x="chocolate", y="winpercent", data=win_num).set_title('Rele
ationship between Win Percent & Chocolate')
```

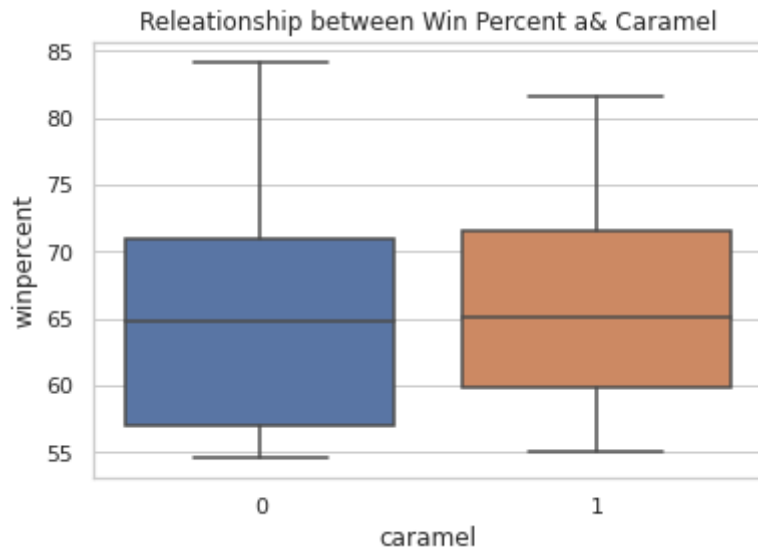
```
Out[45]: Text(0.5, 1.0, 'Releationship between Win Percent & Chocolate')
```



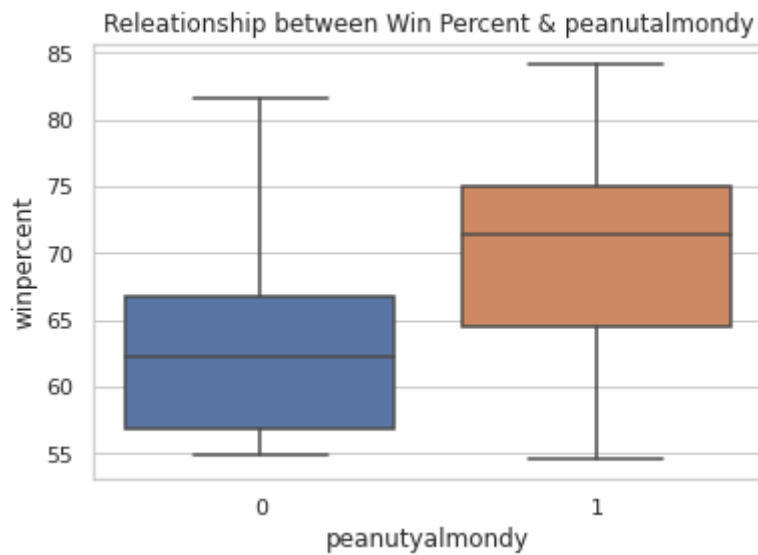
```
In [46]: sns.boxplot(x="fruity", y="winpercent", data=win_num).set_title('Releati
onship between Win Percent & Fruity');
```



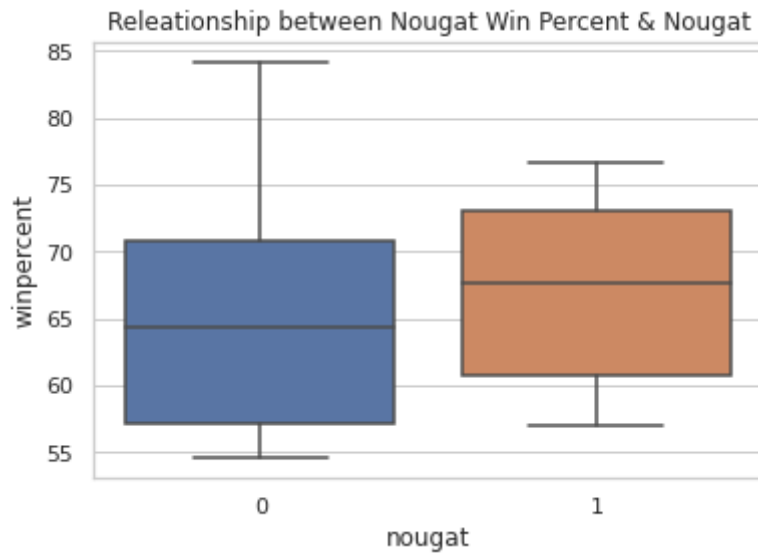
```
In [47]: sns.boxplot(x="caramel", y="winpercent", data=win_num).set_title('Releationship between Win Percent a& Caramel');
```




```
In [48]: sns.boxplot(x="peanutyalmondy", y="winpercent", data=win_num).set_title('Releationship between Win Percent & peanutralmondy');
```



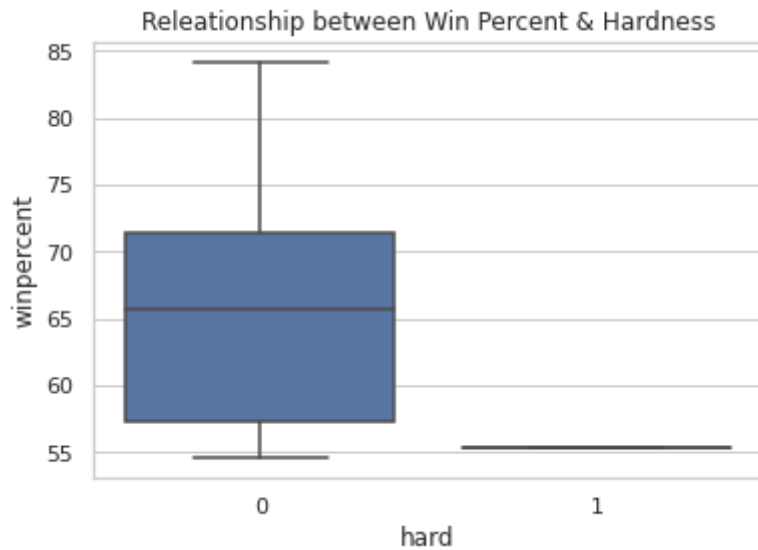
```
In [49]: sns.boxplot(x="nougat", y="winpercent", data=win_num).set_title('Releati  
onship between Nougat Win Percent & Nougat');
```



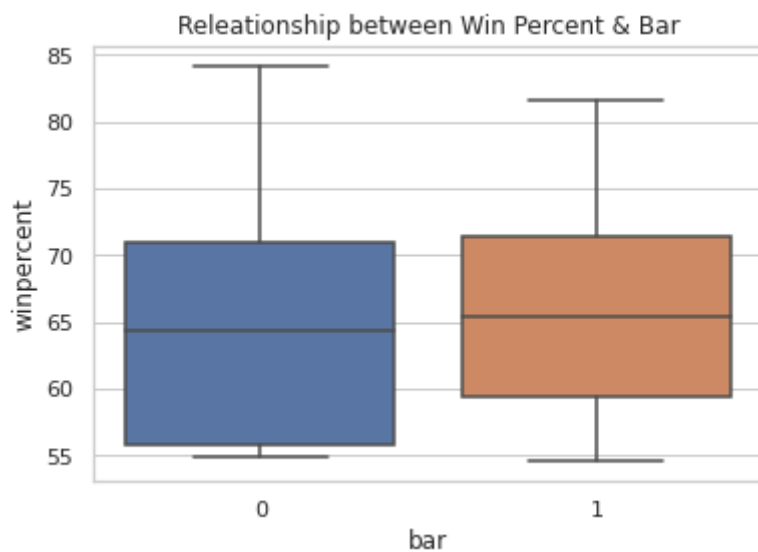
```
In [50]: sns.boxplot(x="crispedricewafer", y="winpercent", data=win_num).set_titl  
e('Releationship between Win Percent & Wafer');
```




```
In [51]: sns.boxplot(x="hard", y="winpercent", data=win_num).set_title('Releation  
ship between Win Percent & Hardness');
```



```
In [52]: sns.boxplot(x="bar", y="winpercent", data=win_num).set_title('Releations  
hip between Win Percent & Bar');
```




```
In [53]: sns.boxplot(x="pluribus", y="winpercent", data=win_num).set_title('Releationship between Win Percent & Pluribus');
```



5) Machine Learning

a) Decision Tree

The decision tree alorgrithm is used here to visually represent feature importance.

The tree has considered all features but is only display with an importance greater then 0.


The decision tree is plotted on a bar chart.

From this bar chart we can see that the most important feature is chocolate and the least is pluribus. Chocolate is in fact significantly more important then the next most important peanutralmondy

Chocolate = 0.772891 vs peanutralmondy = 0.091855

```
In [54]: d_tree = tree.DecisionTreeRegressor(max_depth=3).fit(candy[candy.columns
[1:-1]],candy[candy.columns[-1]])
df_imp = pd.DataFrame.from_dict({'Name':candy.columns[1:-1], 'Importance'
:d_tree.feature_importances_})
df_imp_plt = df_imp.sort_values(by='Importance',ascending=True).reset_in
dex(drop=True)
df_imp_plt[df_imp_plt.Importance>0].plot(kind='barh',x='Name',y='Importa
nce',title='Feature Importance',sort_columns=True,figsize = (15,5),legen
d=False)
```

Out[54]: <matplotlib.axes._subplots.AxesSubplot at 0x7f75815732e8>



```
In [55]: df_imp
```

Out[55]:

	Name	Importance
0	chocolate	0.662891
1	fruity	0.049499
2	caramel	0.079350
3	peanutyalmondy	0.091855
4	nougat	0.000000
5	crispedricewafer	0.000000
6	hard	0.000000
7	bar	0.000000
8	pluribus	0.012091
9	sugarpercent	0.046238
10	pricepercent	0.058076

b) Logistic regression

Logistic Regression is the modelling approach I am taking. It is the bestfitting and least restrictive model.

It describes the relationship between each independent explanatory variables and the dependent binomial response variable.

It will be trying to predict if a candy is chocolatey or not based on all the other features in the dataset

The reason logistic regression is chosen is because it is especially good at predicting boolean values like True or False and has good accuracy for simple datasets.

1. Create a dataframe without the competitor name column.

The competitor name column is a categorical column which will not be used in the machine learning algorithm. It has a unique value in every row so it is impossible to fit a linear model with it. The dataframe is assigned to `x`

2. Create a series with the Chocolate column.

`y` is the responses variable and contains the Chocolate values. Also known as the explanatory variable

```
In [56]: x = candy.drop(['competitorname'], axis = 1)
         y = candy['chocolate']
```

3. Create a test and train set.

`train_test_split` splits the candy set. 75% is allocated to the train set and 25% is allocated to the test set. The train set is a sample of data used to fit the model. The test set is used to provide an unbiased evaluation of the final model fit on the training dataset.

```
In [57]: x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.25
         , random_state=0)
```

4. Perform Logistic Regression.

Assign `logisticRegression()` to `logreg` then fit the model with `logreg.fit()`. The logistic regression is performed on the train set.

```
In [58]: logreg = LogisticRegression()

logreg.fit(X_train,y_train)
```

```
Out[58]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=
True,
                                intercept_scaling=1, l1_ratio=None, max_iter=100,
                                multi_class='auto', n_jobs=None, penalty='l2',
                                random_state=None, solver='lbfgs', tol=0.0001, verbo
se=0,
                                warm_start=False)
```

5. Prediction.

`logreg.predict()` returns a array of the predicted values. The value 1 means the candy is chocolate. The value 0 means the candy is not chocolate. The test set has 22 rows. 8 of the values equal 1 and 14 equal 0. We can test this against the original test set and see which values are different. It appears that the 3rd last value is different in each series.

```
In [59]: y_pred=logreg.predict(X_test)
y_pred
```

```
Out[59]: array([0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 1,
1])
```

```
In [60]: list(X_test['chocolate'])
```

```
Out[60]: [0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 1, 1, 1, 1]
```

6. Confusion Matrix.

It is a performance measurement for machine learning. It will be used for testing the accuracy. It is a table with predicted and actual values.

True Positive = 13 - Predicted positive and is true - predicted its chocolate and its chocolate.

False Positive = 0 - Predicted positive and its false - predicted chocolate and its not.

False Negative = 1 - Predicted negative and its false - predicted not chocolate and its not chocolate.

True Negative = 8 - Predicted negative and its true - predicted not chocolate and it is.

```
In [61]: cnf_matrix = metrics.confusion_matrix(y_test, y_pred)
cnf_matrix
```

```
Out[61]: array([[13,  0],
               [ 1,  8]])
```

7. Accuracy.

The result is a percentage which represents the correct predictions of the test data.

accuracy = True Positive + True Negative/ Total. The result is 95.4545%.

The mean squared error is the average squared difference between the estimated values and true value. It is used to find the percentage error. The result is 0.04545 or 0.45%.

```
In [62]: print("Accuracy:", metrics.accuracy_score(y_test, y_pred))

# or

diagonal_sum = cnf_matrix.trace()
sum_of_all_elements = cnf_matrix.sum()

print("The Accuracy of the model is:", diagonal_sum/sum_of_all_elements*
100)

cost=mean_squared_error(y_test,y_pred)
print("\nThe error is {}".format(cost))
```

Accuracy: 0.9545454545454546

The Accuracy of the model is: 95.45454545454545

The error is 0.04545454545454546

8. Conclusion.

After performing machine learning we can identify that the most important feature is chocolate in candy and we can now predict if a candy is going to be chocolate based on what the other features are at 95.45454545454545% accuracy.